


Tutorial 5: Hadoop MapReduce

1- What is Hadoop?

Hadoop uses a distributed processing architecture called MapReduce in which a task is mapped to a set of servers for processing. The results of the computation performed by those servers are then reduced to a single output set. One node, designated as the master node, controls the distribution of tasks. The following diagram shows a Hadoop cluster with the master node directing a group of slave nodes which process the data.



2-What is Amazon Elastic MapReduce?

With Amazon Elastic MapReduce (Amazon EMR) you can analyze and process vast amounts of data. It does this by distributing the computational work across a cluster of virtual servers running in the Amazon cloud. The cluster is managed using an open-source framework called Hadoop.

Objectives

- Learn how to write MapReduce code using java.
- Learn how to set up a single-note Hadoop deployment and test your MapReduce code.
- Learn how to run MapReduce code in cloud using AWS EMR.

1 Setup a single-node Hadoop Environment:

In this section, we will go over the steps to deploy a single-node Hadoop environment, and run MapReduce applications written in Java using Eclipse.

1. Go to AWS Management Console via AWS Academy portal and create an EC2 instance running Ubuntu Server 18 using the knowledge gathered during **Tutorial 03**.

We will be using this EC2 instance to deploy Hadoop and run MapReduce applications in the subsequent steps.

2. Log in to your EC2 instance via SSH either using the ssh command line utility on MacOS/Linux or Putty on Windows. We learnt how to do this in **Tutorial 03**.
3. In **Tutorial 03**, we installed Java 17 into the EC2 instances that were created. However, this time, we will install Java 1.8 as per the steps provided below. This is because, Hadoop still only supports Java 1.8 fully.

- a. First, check for updates on any package already installed via the following command.

```
ubuntu@ip-172-31-87-252:~$ sudo apt update
```

- b. Then, upgrade all packages that are currently upgradable, as below.

```
ubuntu@ip-172-31-87-252:~$ sudo apt upgrade
```

- c. Finally, we will install Java 1.8, as below.

```
ubuntu@ip-172-31-87-252:~$ sudo apt install openjdk-8-jdk
```

- d. Once done, verify your Java installation as below. You should see the message printed on the command line pointing to Java 1.8.

```
ubuntu@ip-172-31-87-252:~$ java -version
openjdk version "1.8.0_312"
OpenJDK Runtime Environment (build 1.8.0_312-8u312-b07-0ubuntu1~18.04-b07)
OpenJDK 64-Bit Server VM (build 25.312-b07, mixed mode)
```

4. Once Java has been properly installed, download the latest released version of Hadoop version (e.g. 3.3.2) from <https://archive.apache.org/dist/hadoop/core/current/> into your EC2 instance using the steps below.
 - a. Create a directory called **dist** within the home directory (e.g. /home/ubuntu) of the currently logged in user (e.g. ubuntu)

```
ubuntu@ip-172-31-87-252:~$ mkdir dist
```

- b. Change the current working directory to the one created above.

```
ubuntu@ip-172-31-87-252:~$ cd dist
ubuntu@ip-172-31-87-252:~/dist$
```

- c. Verify that you are now in the **dist** directory by running the following command.

```
ubuntu@ip-172-31-87-252:~/dist$ pwd
/home/ubuntu/dist
```

- d. Download the latest released version of Hadoop into your **dist** directory from <https://archive.apache.org/dist/hadoop/core/current/>, as below.

```
ubuntu@ip-172-31-87-252:~/dist$ wget https://archive.apache.org/dist/hadoop/core/current/hadoop-3.3.2.tar.gz
--2022-03-27 11:30:36-- https://archive.apache.org/dist/hadoop/core/current/hadoop-3.3.2.tar.gz
Resolving archive.apache.org (archive.apache.org)... 138.201.131.134, 2a01:4f8:172:2ec5::2
Connecting to archive.apache.org (archive.apache.org)|138.201.131.134|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 639660563 (609M) [application/x-gzip]
Saving to: 'hadoop-3.3.2.tar.gz'

hadoop-3.3.2.tar.gz      100%[=====] 609.07M  16.1MB/s   in 38s

2022-03-27 11:31:15 (15.9 MB/s) - 'hadoop-3.3.2.tar.gz' saved [639660563/639660563]

ubuntu@ip-172-31-87-252:~/dist$
```

5. Extract the content of the .tar.gz file to a location of your choice by running the following command. In this example, we are extracting the Hadoop framework into the dist directory created above.

```
ubuntu@ip-172-31-87-252:~/dist$ tar -xvf hadoop-3.3.2.tar.gz
```

Once properly extracted, you will see the hadoop-3.3.2 directory in the file system, as shown below.

```
ubuntu@ip-172-31-87-252:~/dist$ ls
hadoop-3.3.2  hadoop-3.3.2.tar.gz
ubuntu@ip-172-31-87-252:~/dist$
```

6. In the next step, we will set some environment variables so that we can invoke the hadoop provided tools and scripts on the command line.

For this we will be modifying a hidden file in your home folder of the currently logged in user (e.g. /home/ubuntu/) called .bash_profile, which allows you to set permanent environment variables.

- a. Go to the home folder (e.g. /home/ubuntu) of the currently logged in user (e.g. ubuntu).

```
ubuntu@ip-172-31-87-252:~/dist$ cd /home/ubuntu/
ubuntu@ip-172-31-87-252:~$
```

- b. If the .bash_profile file does not exist in the home folder, just create that file in your home folder using the touch command:

```
ubuntu@ip-172-31-87-252:~$ touch .bash_profile
ubuntu@ip-172-31-87-252:~$
```

- c. You can then open `.bash-profile` in your home folder using a text editor, for example

```
ubuntu@ip-172-31-87-252:~$ vim .bash_profile
ubuntu@ip-172-31-87-252:~$
```

- d. We will do the following tasks:

- Set the `JAVA_HOME` variable to the place where you installed the Java SDK, as well as configuring the `PATH` variable to point to its bin directory.
- Set the `HADOOP_HOME` variable to the place where you extract the Hadoop distribution, as well as configuring the `PATH` variable to point to its bin directory.

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
export PATH=$PATH:$JAVA_HOME/bin
export HADOOP_HOME=/home/ubuntu/dist/hadoop-3.3.2
export PATH=$PATH:$HADOOP_HOME/bin
```

Save the `.bash_profile` file, run the following command to reload your command line to be able to take those variables to take effect.

```
ubuntu@ip-172-31-87-252:~$ source .bash_profile
```

Now you can check and inspect all your environment variables via the command line, as below.

```
ubuntu@ip-172-31-87-252:~/dist$ echo $JAVA_HOME
/usr/lib/jvm/java-8-openjdk-amd64
ubuntu@ip-172-31-87-252:~/dist$ echo $HADOOP_HOME
/home/ubuntu/dist/hadoop-3.3.2
ubuntu@ip-172-31-87-252:~/dist$
```

7. Test your `hadoop version` command; you should be able to see something like:

```
ubuntu@ip-172-31-87-252:~/dist$ hadoop version
Hadoop 3.3.2
Source code repository git@github.com:apache/hadoop.git -r 0bcb014209e219273cb6fd4152df7df713cbac61
Compiled by chao on 2022-02-21T18:39Z
Compiled with protoc 3.7.1
From source with checksum 4b40fff8bb27201ba07b6fa5651217fb
This command was run using /home/ubuntu/dist/hadoop-3.3.2/share/hadoop/common/hadoop-common-3.3.2.jar
ubuntu@ip-172-31-87-252:~/dist$
```

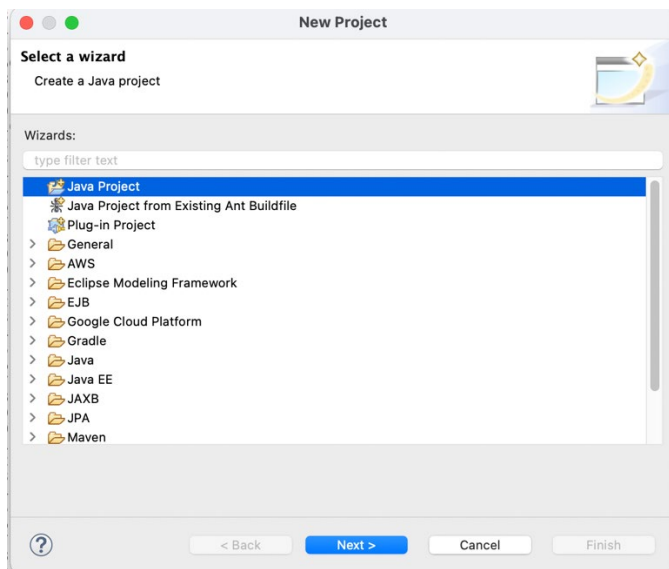
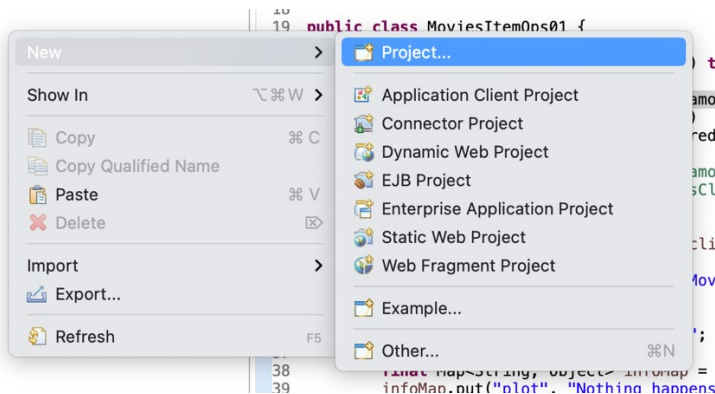
Now we are ready to run our MapReduce applications. In summary, we will write code and compile our MapReduce code on Eclipse. Then, we will export our MapReduce applications as a jar file and run it on top of Hadoop using the command line.

1.1 Set-up Eclipse to write MapReduce applications

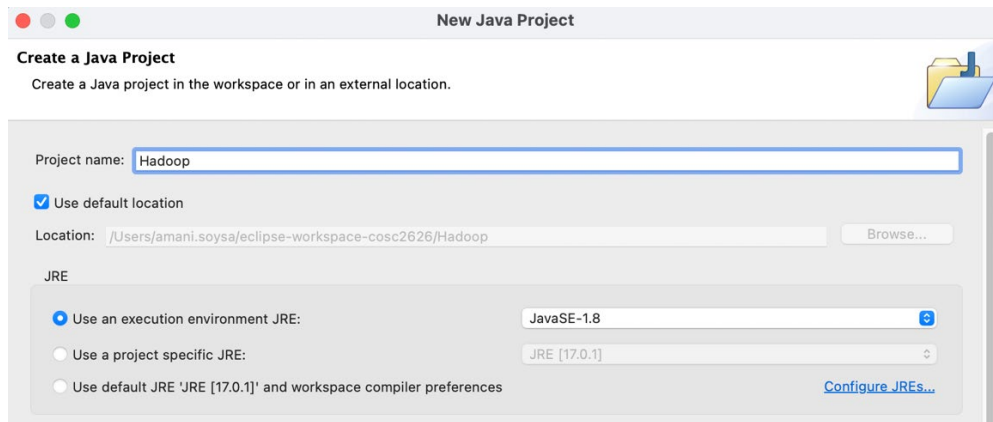
1. Download the same Hadoop distribution (i.e. hadoop-3.3.2.tar.gz) from <https://archive.apache.org/dist/hadoop/core/current/>, now into **your local development environment**. Extract it to your preferred location in the file system.

This Hadoop distribution not only bundles the Hadoop distribution that runs MapReduce applications, but also the client libraries required to write the aforementioned applications. These client libraries could be located in `hadoop-<version>/lib` directory.

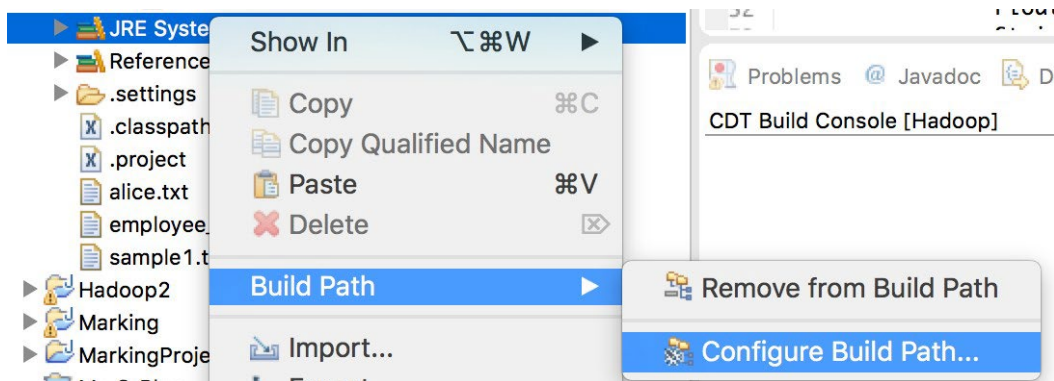
2. Create a Java project on Eclipse called **Hadoop**. To do this, either right click on the **Project Explorer** or select **File -> New -> Java Project**.



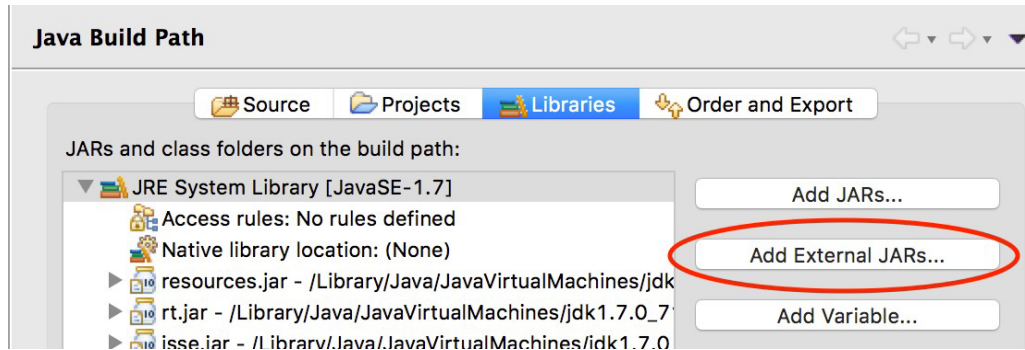
While creating your Java project, make sure you select **Java 1.8** as the JRE associated with the project, as shown below.



3. Now, we need to add 2 more external jar files to the classpath of our newly created **Hadoop** project. These jar files correspond to the client libraries required to write MapReduce applications, which we previously downloaded as part of the Hadoop distribution.
 - a. Right click on the JRE system library of your project, select Build Path, **Configure Build Path**, as below.



- b. Click on **Add External JARs**.

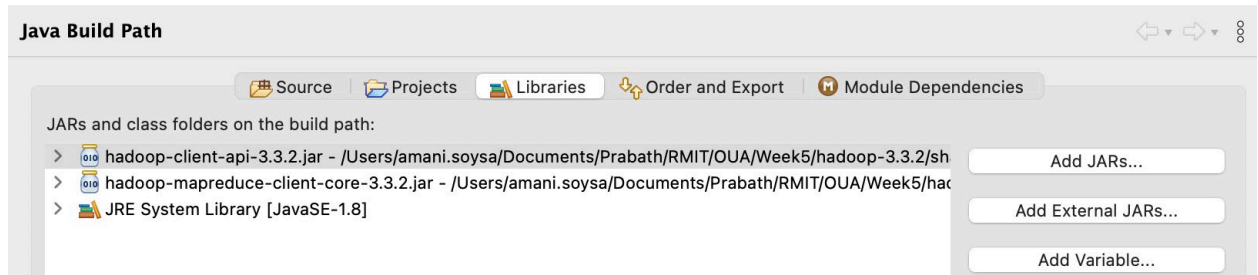


- c. Search in the folder where you have extracted Hadoop and add the following 2 jar files (Shown below is the full path to those jar files).

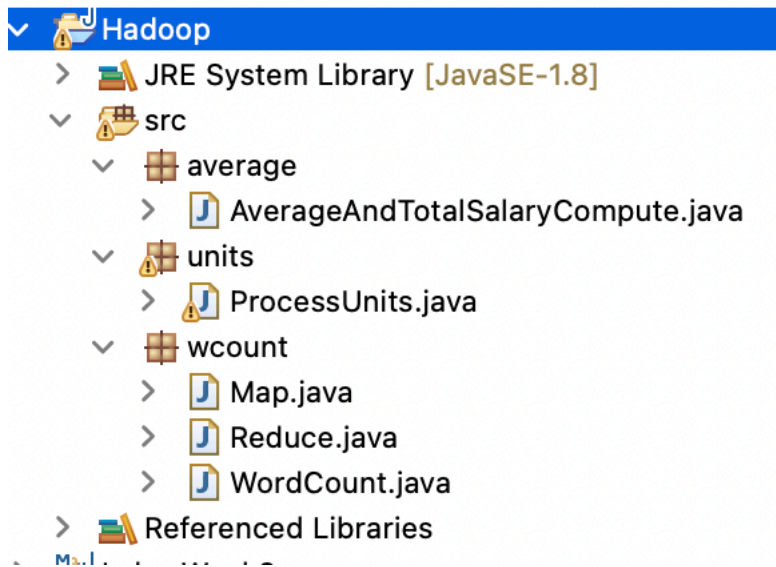
Note: Here, **HADOOP_HOME** refers to the folder to which you have downloaded the **hadoop-<version>.tar.gz** package.

```
$HADOOP_HOME/share/hadoop-3.3.2/client/hadoop-client-api-3.3.2.jar
$HADOOP_HOME/share/hadoop-3.3.2/mapreduce/hadoop-mapreduce-client-core-3.3.2.jar
```

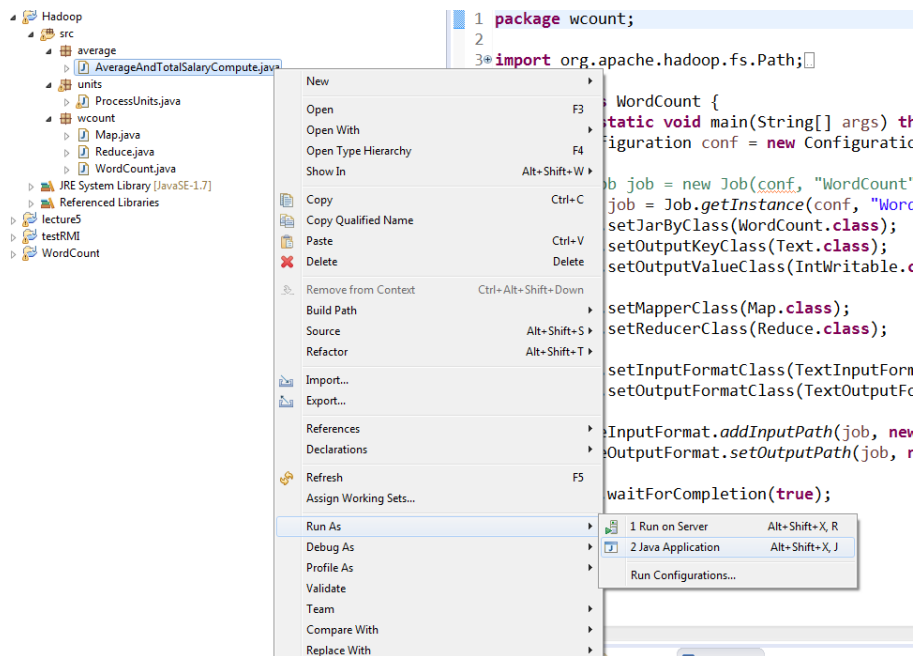
- d. After you have added those 2 jars file, you should see something like following.



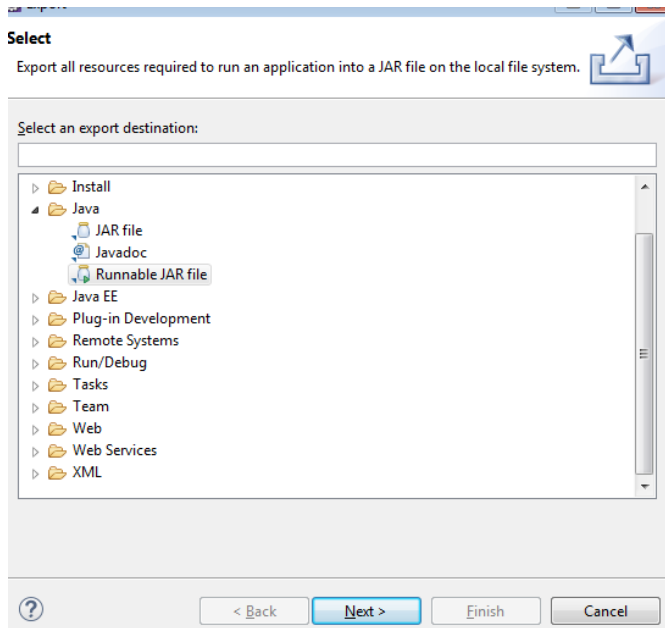
4. In the src folder of the Java project you created (e.g. Hadoop), add 3 new packages named
 - 1) average
 - 2) units
 - 3) wcount
5. Now download the **mrcode.zip** package from the Canvas via *Modules -> Week 5: additional learning materials* section. Extract this package and copy the **Java class files** inside that package into your **Hadoop** project according to the following structure.



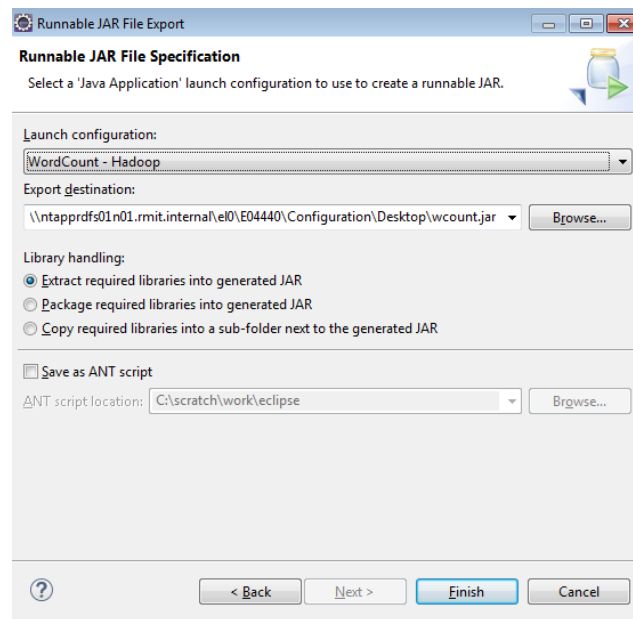
6. Closely examine the implementation of the sample code you copied into your project. Try to understand how the MapReduce paradigm is implemented in these sample applications. Now run each of 3 packages as a **Java application**, as below. Don't worry if you get exceptions.



7. Now, export each of the aforementioned 3 packages as runnable jars. For example, to export wcount package as a jar, right click on the package name and select export. In the Export window, select **Java > Runnable jar file**.



8. In next window, Select the main file as Launch configuration. For example, for wcount package the main function is in WordCount.java file. Also give the jar file name same as package name (e.g. wcount.jar) as in the following screenshot.



9. Click **Finish**. Click Ok if you see any warning. Now your jar files are ready to be run on Hadoop.
10. Now, create a new directory called **jars** inside the home directory (e.g. /home/ubuntu/) of the logged in user (e.g. ubuntu) within the EC2 instance we

created in a previous step. Copy all generated runnable jar files you have created earlier into this **jars** directory.

Also create another folder called **inputs** inside home directory of the logged in user within your EC2 instance. Now, copy all the text files in **mrcline/inputs** directory you downloaded from Canvas into this directory.

The folder hierarchy within the home directory (e.g. /home/ubuntu/) of the logged in user (e.g. ubuntu) should now look like the following.

```
ubuntu@ip-172-31-87-252:~$ ls
dist  inputs  jars
ubuntu@ip-172-31-87-252:~$
```

Run the first project: Electrical Consumption

Given below is the data regarding the electrical consumption of an organization. It contains the monthly electrical consumption and the annual average for various years.

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec	Avg
1979	23	23	2	43	24	25	26	26	26	26	25	26	25
1980	26	27	28	28	28	30	31	31	31	30	30	30	29
1981	31	32	32	32	33	34	35	36	36	34	34	34	34
1984	39	38	39	39	39	41	42	43	40	39	38	38	40
1985	38	39	39	39	39	41	41	41	00	40	39	39	45

If the above data is given as input, we have to write applications to process it and produce results such as finding the year of average usage that is greater than 30.

Connect to your EC2 instance via SSH. Then run the following command from the home

directory (e.g. /home/ubuntu) of the logged in user (e.g. ubuntu).

```
hadoop jar jars/units.jar inputs/sample1.txt output-units
```

Here, we used the hadoop command line utility (CLI) to run the MapReduce application embedded in **units.jar** file. Hadoop then processes the data inside the **input/sample1.txt** file and generate the output inside **output-units** folder. Note that here you don't need to create an output folder. Hadoop will automatically create it and place the result in it.

After Hadoop finish running you will see output like the following

```
2022-03-28 07:44:19,628 INFO mapred.LocalJobRunner: Finishing task: attempt_local132865647_0001_r_000000_0
2022-03-28 07:44:19,628 INFO mapred.LocalJobRunner: reduce task executor complete.
2022-03-28 07:44:19,971 INFO mapreduce.Job: Job job_local132865647_0001 running in uber mode : false
2022-03-28 07:44:19,972 INFO mapreduce.Job: map 100% reduce 100%
2022-03-28 07:44:19,973 INFO mapreduce.Job: Job job_local132865647_0001 completed successfully
2022-03-28 07:44:19,986 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=38923154
    FILE: Number of bytes written=40496271
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
  Map-Reduce Framework
    Map input records=5
    Map output records=5
    Map output bytes=45
    Map output materialized bytes=39
    Input split bytes=88
    Combine input records=5
    Combine output records=3
    Reduce input groups=3
    Reduce shuffle bytes=39
    Reduce input records=3
    Reduce output records=3
    Spilled Records=6
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=38
    Total committed heap usage (bytes)=243113984
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=343
  File Output Format Counters
    Bytes Written=36
ubuntu@ip-172-31-87-252:~$
```

Let's now look at the output.

```
ubuntu@ip-172-31-87-252:~$ cd output-units/
ubuntu@ip-172-31-87-252:~/output-units$ ls
_SUCCESS part-00000
ubuntu@ip-172-31-87-252:~/output-units$ cat part-00000
1981    34
1984    40
1985    45
ubuntu@ip-172-31-87-252:~/output-units$
```

When Hadoop has finished running the MapReduce application successfully, the content of the output folder looks like what is shown above. To view the content of the output file, you use the **cat** command.

Run the second project: Average Salary by Genders

Let's now run another application, which outputs the average salary of employees grouped by their gender. To do this, run the second runnable jar file named **average.jar** using the following command.

```
hadoop jar jars/average.jar inputs/employee_records.txt
output-average
```

Upon successful execution of the aforementioned, command your output will look like what is shown below.

```
ubuntu@ip-172-31-87-252:~$ cd output-average/
ubuntu@ip-172-31-87-252:~/output-average$ ls
_SUCCESS part-r-00000
ubuntu@ip-172-31-87-252:~/output-average$ cat part-r-00000
F      Total: 291800.0 :: Average: 7117.073
M      Total: 424363.34 :: Average: 6333.7812
ubuntu@ip-172-31-87-252:~/output-average$
```

Run the third project: Word Count

Once again remove the employee_records.txt from input folder and remove the output folder. Then copy alice.txt in input folder. Now test the third package **wcount** by using the command.

```
hadoop jar jars/wcount.jar inputs/alice.txt output-wcount
```

Upon successful execution of the aforementioned, command your output will look like what is shown below.

```

ubuntu@ip-172-31-87-252:~$ cd output-wcount/
ubuntu@ip-172-31-87-252:~/output-wcount$ ls
_SUCCESS part-r-00000
ubuntu@ip-172-31-87-252:~/output-wcount$ cat part-r-00000
a      269
about  32
above  1
accustomed  1
across  5
actually  1
added   6
addressed  1
adoption  1
adventures  3
advice  1
advisable  1
advise  1
affectionately  2
afore  1
afraid  4
after  16
again  7
against  5
agony  1
air  1
alas  1
alice  103
all  52
almost  2
along  4
already 1
alternately  1
always  2
am  4
among  9
an  18
and  312
anger  1
angry  3
animal 1
animals 4
ann  1
another 9
answer 2
anxiously  5
any  13
anyone 2
anything  4
anywhere  1
appearance  1
archbishop  2
are  13
arm  7
arms  4
around 1
as  86
asked  3

```

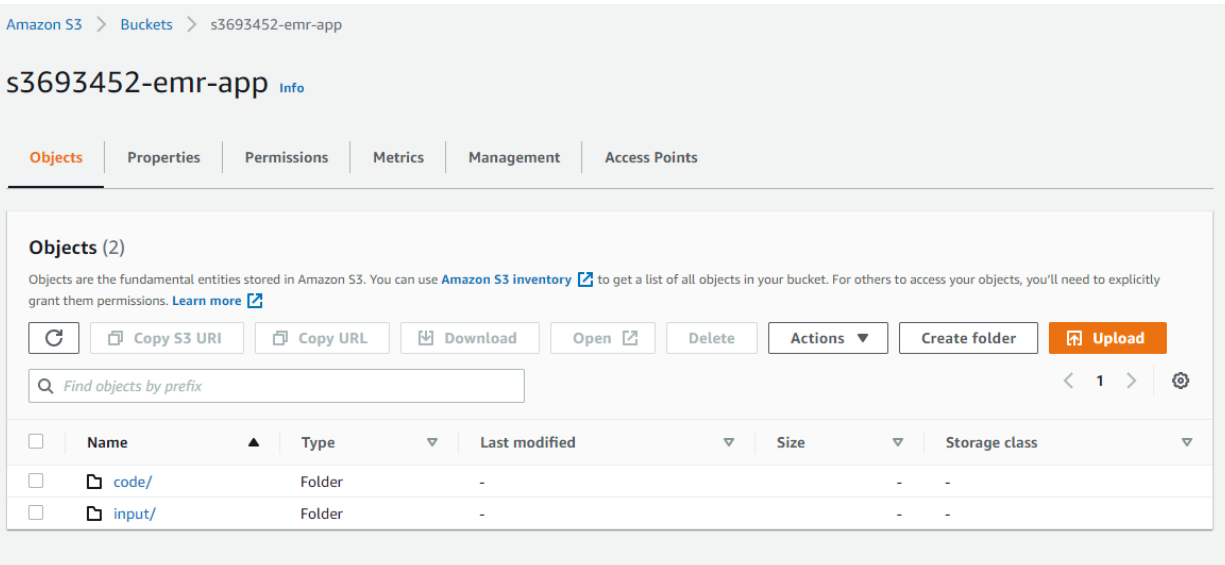
Here, Hadoop is generating only one output file because we are using a single node cluster. That means, there is only one processing node at work. This process is helpful when your dataset is small and when you need to test your developed MapReduce code locally. But if your dataset is large a single node cluster will not be enough, and you need to take the advantage of cloud by using a Multi-Node cluster. AWS has such support in EMR. Now we will use our developed MapReduce code in EMR.

2 Solve word count problem using EMR (Amazon Elastic MapReduce)

2.1 Prepare your scripts and input data

1- Log in to your AWS management console via AWS Academy portal. Select AWS S3 and create a bucket named **sxxxxxx-emr-app** (sxxxxxx refers to RMIT your student id).

2- Create two folders inside your bucket and name them **code** and **input**.



3- Inside **code** folder, upload the **wcount.jar** file you created before.

4- Download the text file from here:

<http://www.umich.edu/~umfandsf/other/ebooks/alice30.txt> and upload alice30.txt into the **input** folder of your Bucket.

2.2 Create Amazon MapReduce cluster

- 1- Login to AWS management console and from services select "EMR".
- 2- Click on **Create cluster** and you will see a screen like following.

General Configuration

Cluster name

☒ Logging ⓘ

S3 folder ⓘ

Launch mode ☒ Cluster ⓘ ☐ Step execution ⓘ

Software configuration

Vendor ☒ Amazon ☐ MapR

Release ⓘ

Applications

- ☒ Core Hadoop: Hadoop 2.7.2 with Ganglia 3.7.2, Hive 2.1.0, Hue 3.10.0, Mahout 0.12.2, Pig 0.16.0, and Tez 0.8.4
- ☐ HBase: HBase 1.2.2 with Ganglia 3.7.2, Hadoop 2.7.2, Hive 2.1.0, Hue 3.10.0, Phoenix 4.7.0, and ZooKeeper 3.4.8
- ☐ Presto: Presto 0.150 with Hadoop 2.7.2 HDFS and Hive 2.1.0 Metastore
- ☐ Spark: Spark 2.0.0 on Hadoop 2.7.2 YARN with Ganglia 3.7.2 and Zeppelin 0.6.1

- 3- Give a name to your cluster. In the **Logging S3 folder location** field browse and select s3://sxxxxxxx-emr-app, where sxxxxxxx your RMIT student id.
- 3- In Launch mode select **Step execution**. Then, select step type as **Custom JAR** as in the screenshot below, and click **Configure**.

General Configuration

Cluster name

☒ Logging ⓘ

S3 folder

Launch mode ☐ Cluster ⓘ ☒ Step execution ⓘ

Add steps

A step is a unit of work submitted to an application running on your EMR cluster. EMR programmatically installs the applications needed to execute the added steps. [Learn more](#)

Step type

4- A popup window will open. Fill it as shown in the following figure. In the arguments you have to provide the input location as the first arguments and the output location as the second arguments as shown here.

s3://sxxxxxx-emr-app/input/ s3://sxxxxxx-emr-app/output/

Add step

Step type Custom JAR

Name* Custom JAR

JAR location* s3://s3693452-emr-app/code/wcount.jar

Arguments s3://s3693452-emr-app/input
s3://s3693452-emr-app/output

Action on failure Continue

What happens if the step fails

Cancel Add

JAR location maybe a path into S3 or a fully qualified java class in the classpath.

These are passed to the main function in the JAR. If the JAR does not specify a main class in its manifest file you can specify another class name as the first argument.

Also select **Terminate cluster** as the **Action on failure**.

Note that, in case of Output **S3 location**, you don't need to create **output** directory explicitly in your bucket. The EMR service will do it for you. However, make sure you type '/output/' after your bucket name in the aforementioned configuration.

Also note that, in this cluster you are using 3 m3.xlarge VM instances. And minimum number of instance also should be 3. But, you can use more powerful instance types and higher number of instances here, which would improve the throughput as well as

scalability depending on the requirements of the application. The more resources you use the more you will need to pay. We kept everything at a minimum for this exercise. However, keep in mind that more powerful instances and a large number of instances in your Hadoop cluster will improve the performance of MapReduce applications.

5- Keep **Security and Access** as default.

9-Now click on **Create cluster** and wait until it finishes. It can take 10 minutes or more to complete. It will show "Starting" for some time. It may take more than 5-10 minutes to start. **Starting** means the system is configuring the clusters you have created.

Cluster: My cluster Starting

Summary

Application user interfaces

Monitoring

Hardware

Configurations

Events

Steps

Bootstrap actions

Summary

Configuration details

Application user interfaces

Network and hardware

Security and access

ID: j-44RHRO9SQFCJ

Creation date: 2022-03-28 19:19 (UTC+11)

Elapsed time: 1 second

After last step completes: Cluster auto-terminates

Termination protection: Off [Change](#)

Tags: -- [View All / Edit](#)

Master public DNS: --

Release label: emr-5.35.0

Hadoop distribution: Amazon 2.10.1

Applications: --

Log URI: s3://s3693452-emr-app/

EMRFS consistent view: Disabled

Custom AMI ID: --

Persistent user interfaces : --

On-cluster user interfaces : --

Availability zone: --

Subnet ID: [subnet-07b4f3e77305888a3](#)

Master: Provisioning 1 m3.xlarge

Core: Provisioning 2 m3.xlarge

Task: --

Cluster scaling: Not enabled

Auto-termination: Terminate if idle for 1 hour

Key name: --

EC2 instance profile: EMR_EC2_DefaultRole

EMR role: EMR_DefaultRole

Visible to all users: All [Change](#)

Security groups for Master:

Security groups for Core & Task:

Amazon EMR

EMR Studio

EMR on EC2

Clusters

Notebooks

Git repositories

Security configurations

Block public access

VPC subnets

Events

EMR on EKS

Virtual clusters

Create cluster

View details

Clone

Terminate

Filter: All clusters

1 cluster (all loaded)

	Name	ID	Status	Creation time (UTC+11)	Elapsed time	Normalized instance hours
<input type="checkbox"/>	My cluster	j-44RHRO9SQFCJ	Starting	2022-03-28 19:19 (UTC+11)	1 minute	0

Summary

Steps

Bootstrap actions

Master public DNS: --

Termination protection: Off [Change](#)

Tags: -- [View All / Edit](#)

Master: Provisioning 1 m3.xlarge

Core: Provisioning 2 m3.xlarge

Task: --

[View cluster details](#)

[View monitoring details](#)

[Add step](#)

[View all interactive jobs](#)

Name	Status	Start time (UTC+11)	Elapsed time
Custom JAR	Pending	--	--
Setup hadoop debugging	Pending	--	--

No bootstrap actions available

When MapReduce job will start the status will show "Running" like following screen.

The screenshot shows the Amazon EMR console interface. On the left is a navigation menu with options like EMR Studio, EMR on EC2, Clusters, Notebooks, Git repositories, Security configurations, Block public access, VPC subnets, Events, EMR on EKS, and Virtual clusters. The main panel displays a list of clusters with a filter set to 'All clusters'. One cluster, 'My cluster' (ID: j-44RHRO9SQFCJ), is shown with a status of 'Running'. Below the cluster list, there are sections for Summary, Steps, and Hardware. The Summary section shows Master public DNS, Termination protection, and Tags. The Steps section shows two steps: 'Custom JAR' and 'Setup hadoop debugging', both with a status of 'Pending'. The Hardware section shows Master and Core instances in 'Running' status.

Wait for some more time. If the status shows "Terminated steps completed" like following screen, it means MapReduce job has finished and the outputs are now ready for viewing.

The screenshot shows the Amazon EMR console interface. The cluster 'My cluster' (ID: j-44RHRO9SQFCJ) now has a status of 'Terminated steps completed'. The Summary section shows the Master public DNS as 'ec2-34-207-243-247.compute-1.amazonaws.com'. The Steps section shows two steps: 'Custom JAR' and 'Setup hadoop debugging', both with a status of 'Completed'. The Hardware section shows Master and Core instances in 'Terminating' status.

2.3 View the output result

Once your job is marked as completed, go to AWS Management Console via AWS Academy portal, Select **AWS S3**, and click on the **output** folder you specified. You will see some files named part-0000, part-0001,.. which contain the results of your MapReduce job. The number of output files generated depends on the **number of reducers** at work to complete a given MapReduce application. This is managed by EMR.

Amazon S3 > Buckets > s3693452-emr-app > output/

output/ Copy S3 URI

Objects Properties

Objects (8)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Refresh Copy S3 URI Copy URL Download Open Delete Actions Create folder Upload

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	_SUCCESS	-	March 28, 2022, 19:27:29 (UTC+11:00)	0 B	Standard
<input type="checkbox"/>	part-r-00000	-	March 28, 2022, 19:27:17 (UTC+11:00)	2.7 KB	Standard
<input type="checkbox"/>	part-r-00001	-	March 28, 2022, 19:27:17 (UTC+11:00)	2.5 KB	Standard
<input type="checkbox"/>	part-r-00002	-	March 28, 2022, 19:27:28 (UTC+11:00)	2.4 KB	Standard
<input type="checkbox"/>	part-r-00003	-	March 28, 2022, 19:27:27 (UTC+11:00)	2.6 KB	Standard
<input type="checkbox"/>	part-r-00004	-	March 28, 2022, 19:27:17 (UTC+11:00)	2.5 KB	Standard
<input type="checkbox"/>	part-r-00005	-	March 28, 2022, 19:27:27 (UTC+11:00)	2.7 KB	Standard
<input type="checkbox"/>	part-r-00006	-	March 28, 2022, 19:27:27 (UTC+11:00)	2.6 KB	Standard

You can download all the files and open them in a text editor to view the results. The results contain the word and number of occurrences of that word in "alice30.txt" file, as below.

part-r-00000 - Notepad

File	Edit	Format	View	Help
acceptance		1		
added	20			
addressed		2		
adventures		4		
advisable		1		
affectionately		1		
alone	2			
altogether		1		
and	766			
angry	3			
any	36			
arm	7			
atom	2			
attempts		1		
back	29			
baked	1			
barking	1			
beautify		1		
because	12			
before	19			
began	47			
belongs	2			
blades	1			
blow	2			
bones	1			
bother	1			
bottom	4			
branches		1		
buttercup		1		
called	15			
canvas	1			
carrying		2		
caterpillar		11		
centre	1			
certain	2			

Important note: The commands we run to launch our MapReduce job in this example are set to automatically terminate the EC2 instances. However, as you get into more custom commands, this will not necessarily be the case. In those situations, you **HAVE TO** terminate the instances yourself. Otherwise, you will be charged for the time these instances are running (even if the job completed). To terminate an instance, you can Go on Amazon Management Console via AWS Academy portal → Elastic Map Reduce, select your Job Flow

and click on Terminate.